

Mark Shepard, Research Statement

September 2025

Overview

My research studies the **economics of health insurance**, with a focus on the U.S. health insurance system. Within economics, my work bridges health economics, industrial organization, and public economics. I also contribute to the health policy literature with targeted insights on policy-relevant issues.

Health systems around the world face persistent challenges in managing costs, ensuring access to care, and incentivizing quality and innovation. Nowhere are these challenges more acute than in the United States, which spends over \$5 trillion annually – roughly 18% of GDP – while still leaving 26 million people without health insurance coverage. At the center of these pressures is health insurance: it is the primary way modern health systems finance care, yet also one of the most complex and difficult institutions to design. My research aims to make sense of the complex economic forces that shape health insurance and to evaluate how reforms might make it work better.

A central theme of my research is the need to understand *health insurance as a system*, not merely a patchwork of markets and programs. A true insurance system sets baseline, non-negotiable goals—such as universal access to essential care and limits on overall budgets—and designs structures to guarantee these goals are met. Markets, by contrast, rely on decentralized, voluntary trades that, while individually beneficial, often neglect broader impacts. The United States, shaped by a history of incremental reforms, lacks a coherent health insurance system: its fragmented programs create conflicting incentives, fail to control costs, and leave millions uninsured. My work shows the difficulty of achieving meaningful progress without reforms addressing this underlying systems problem in U.S. health insurance.¹

This systems perspective plays out in two core areas of my research:

1. **Challenge of Universal Coverage:** My work investigates why uninsurance remains stubbornly high, despite decades of reforms aimed at expanding coverage. A central finding is the *large impact of seemingly small enrollment frictions* — such as modest premiums and administrative hassles — especially for younger, healthier, and lower-income individuals who policymakers most want to reach. These findings suggest that simply making insurance available and affordable – the main strategy of the Affordable Care Act and prior reforms – is likely to fall well short of universal coverage. Instead, achieving this goal requires organizing the system across programs to remove frictions and *automatically* enroll people in coverage for which they are eligible.
2. **Challenge of Insurance Competition:** Insurance systems often embed markets – as in market-based national insurance systems in countries like the Netherlands and Germany, and market-based U.S. programs like the ACA and Medicaid managed care. In America alone, these markets now cover over 100 million people and account for over \$1 trillion in public spending. My research addresses practical policy questions—such as how to design subsidies or how to balance tradeoffs in risk adjustment—while also offering deeper insights into the classic problem of *adverse selection*. Across multiple papers, I find that adverse selection has a more pervasive impact than

¹ I lay out this broad case in a 2023 paper in the *Journal of Economic Perspectives*, “[Achieving Universal Health Insurance Coverage in the United States: Addressing Market Failures or Providing a Social Floor?](#)” (with Kate Baicker and Amitabh Chandra), which I discuss further below.

previously realized because it pushes insurers to distort their products to avoid the sick and attract the healthy. For example, I show how selection pushes insurers towards “narrow” network plans that limit costs by excluding the most expensive and prestigious hospitals. My work suggests new policy insights to help manage the tradeoffs involved.

In what follows, I summarize my specific papers under these two topic headings.

Topic #1: Challenge of Universal Coverage

The U.S. has long struggled to cover the uninsured. Although straightforward in principle, that goal has proved challenging, despite massive programs to provide subsidized health insurance. While uninsured rates fell immediately after the ACA’s enactment, they have been stuck at around 9-11% in the years since 2016. Why have gains not been larger, and what policies could make a difference?

In “[Subsidizing Health Insurance for Low-Income Adults: Evidence from Massachusetts](#)” (with Amy Finkelstein and Nathaniel Hendren, *AER* 2019), we examine how modest financial prices (premiums) affect health insurance take-up among the poor. Using a natural experiment from subsidy discontinuities in Massachusetts, we find that even small premiums significantly reduce enrollment—each \$40 monthly increase reduces enrollment by 25%. Reductions are larger among the young and healthy (consistent with adverse selection), indicating that cutting subsidies makes insurance risk pools sicker. These results provide some of the best-identified evidence of a sharp tradeoff between subsidy size and health insurance enrollment – a finding of perennial relevance in policy debates.² Further, we challenge economists’ conventional understanding of the rationale for subsidies based on adverse selection alone. We find that even correcting for selection, most low-income people value insurance far below its marginal cost. Instead, the case for subsidizing insurance likely needs to rest on either behavioral biases (people undervaluing insurance) or savings on “uncompensated care” incurred when the uninsured get sick.

If financial costs are one barrier to covering the uninsured, a second barrier is the hassles and complexities involved with insurance enrollment. In “[Do Ordeals Work for Selection Markets? Evidence from Health Insurance Auto-Enrollment](#)” (with Myles Wagner, *AER* 2025), we show that hassles significantly impede insurance enrollment, even when coverage is free. Leveraging the removal of an auto-enrollment policy in Massachusetts – a change that added a step to the take-up process – we find that this modest “ordeal” reduces enrollment by 33% and differentially excluded young, healthy, and economically disadvantaged people. These findings have two main implications. First, they deepen our understanding of the insurance take-up challenge, pointing to the large effect of administrative frictions, not just affordability. This “*administrative burden*” framing of health insurance coverage (Herd & Moynihan, 2018) is a fruitful avenue for future research given the complex and fragmented nature of the U.S. health insurance system. Second, our results contribute to the classic debate on when “ordeals” work well in welfare programs (Nichols & Zeckhauser, 1982). We show both theoretically and empirically how adverse selection (a classic insurance issue) tends to undermine the classic case for ordeals targeting.

A few shorter papers contribute to this broader theme of the challenge of health insurance coverage:

² For example, we released our paper in 2017 amidst debate about ACA “repeal and replace,” drawing coverage from [David Leonhardt in the NY Times](#). Our results continue to resonate in ongoing debates about subsidy increases vs. cuts and have been used by the Congressional Budget Office as a key source of evidence.

- In [“Small Marketplace Premiums Pose Financial and Administrative Burdens”](#) (with **Adrianna McIntyre and Tim Layton**, *Health Affairs* 2024), we show how small premium increases in ACA markets – from \$0 to nominal amounts (less than \$10 per month) – lead sizable shares of enrollees to drop out of coverage, and provide evidence that this attrition is explained by the administrative friction of initiating premium payment, rather than unaffordability.
- In [“Can Automatic Retention Improve Health Insurance Market Outcomes?”](#) (with **Adrianna McIntyre and Myles Wagner**, *AEA P&P* 2021), we show how an “automatic retention” policy used by Massachusetts prior to the ACA effectively smoothed this friction by automatically switching people who lapse on premiums to a free plan if available.
- In a related paper ([“Turnover in Zero-Premium Status Among Health Insurance Marketplace Plans Available to Low-Income Enrollees”](#), with **Ed Kong and Adrianna McIntyre**, *JAMA Health Forum* 2022), we show that this zero-to-positive premium transition is quite common in ACA markets.
- Adrianna McIntyre and I summarize policy take-aways from this line of research in a **2022 perspectives paper for the *New England Journal of Medicine*** ([“Automatic insurance policies – important tools for preventing coverage loss”](#)).

An overall theme from my research is that incremental reforms to make health insurance available and affordable – the main approach of the ACA and decades of U.S. reforms – is likely to fall well short of universal coverage. Instead, this goal requires a system of universal coverage in which health insurance is automatic and free, aside from any tax-based contributions. I lay out these ideas in a **2023 *Journal of Economic Perspectives* paper**, [“Achieving Universal Health Insurance Coverage in the United States: Addressing Market Failures or Providing a Social Floor?”](#) (with **Kate Baicker and Amitabh Chandra**). In this research synthesis, we argue for a “social floor” approach to universal coverage (as opposed to the U.S.’s piecemeal approach) and highlight the role economic analysis has to play in helping design and make cost-quality tradeoffs within a universal coverage system.

Topic #2: Challenge of Insurance Competition

Many health insurance systems are centered around regulated markets, a model sometimes called “managed competition.” In the U.S., key examples include the ACA marketplaces and Medicaid managed care;³ internationally, they include market-based systems like the Netherlands, Switzerland, Germany, Israel, Chile, and Colombia. These systems face an important challenge: designing regulations and incentives in a way that steers market competition toward desirable outcomes, rather than market failures. How well this is occurring – and how it could be done better – is a key focus of my research.

In [“Hospital Network Competition and Adverse Selection: Evidence from the Massachusetts Health Insurance Exchange”](#) (*American Economic Review*, 2022), I provide evidence on the striking phenomenon in the ACA of proliferation of plans that cover only a very limited or “narrow” network of hospitals and doctors. Using both a natural experiment and a structural model, this paper provides the first direct evidence that adverse selection pushes insurers towards narrower networks – particularly, by

³ I co-authored two 2018 book chapters summarizing the market design and competition issues involved with the [ACA Marketplaces](#) and [Medicaid managed care](#). While some of my earlier work, these chapters contain many insights that continue to shape my thinking and research agenda on health insurance markets.

excluding the top “star” hospitals to induce their highest-cost consumers to switch away. Adverse selection occurs through a novel channel. Plans covering star hospitals attract not just *sicker* people – the typical mechanism for adverse selection – but also people who incur high costs because they *prefer using the expensive star hospital*. The latter “selection on preferences” channel presents a challenge to standard policy remedies like risk adjustment and connects adverse selection to fundamental tradeoffs involved with cost, quality, and high hospital prices due to market power.

In “[Adverse selection and network design under regulated plan prices: Evidence from Medicaid](#)” (with Amanda Kreider, Tim Layton, and Jacob Wallace; *Journal of Health Economics*, 2024), we show that similar adverse selection disincentives arise for covering cancer hospitals in Medicaid. This disincentive occurs despite the absence of premiums in Medicaid (all plans are \$0), which is interesting because price competition is the standard channel for risk selection in economic models. We show that fixing adverse selection through the standard policy (risk adjustment) is likely to be quite hard and instead point to a simple alternative – explicit “quality” bonuses for plans that cover the cancer hospital.

In addition to competition on quality, price competition is an important goal of the ACA. In “[Price-Linked Subsidies and Imperfect Competition in Health Insurance](#)” (with Sonia Jaffe, *AEJ: Economic Policy*, 2020), we study how subsidy design influences price competition. Subsidies in the ACA (and many other programs) have the interesting feature that they are *linked to prices* set by insurers – if prices rise, so do subsidies. This “price-linked” subsidy design is intended to ensure affordability by protecting low-income consumers from bearing premium increases. But we show that it has an unintended consequence of *weakening insurer price competition*. Using both theory and a structural model, we find that the policy raises prices modestly (about 1-6%) but can have important advantages in the face of substantial uncertainty or political constraints. We argue that price-linked subsidies help explain why ACA insurance markets did not unravel in the face of major shocks during the first Trump administration: even though (gross) premiums rose, subsidies rose in tandem and prevented the enrollee exodus that creates a death spiral.

“[The Two Margin Problem in Insurance Markets](#)” (with Michael Geruso, Tim Layton, and Grace McCormack, *Review of Economics and Statistics*, 2023), studies an issue related to both enrollment and competition. We point out that adverse selection occurs in two different ways in insurance markets (or along two “margins”) that are typically treated separately but may interact. Within the market (the intensive margin), adverse selection makes it hard for *high-quality* plans to survive. But on the extensive margin (in vs. out of the market), adverse selection leads to higher prices and *lower enrollment* in insurance. We develop a theoretical framework that incorporates both margins of selection and show that policies intended to address one margin often have unintended consequences for the other. For instance, when mandates bring healthier people into the market (extensive margin), these people tend to choose lower-quality, low-price plans. This worsens adverse selection between high- vs. low-quality plans (intensive margin). We find that selection policy should adjust in a complementary way: stronger mandates imply stronger policies to ensure minimum quality; weaker mandates imply weaker quality regulation.

In “[Adverse Selection and \(un\)Natural Monopoly in Insurance Markets](#)” (with Timothy Layton and Edward Kong; *Reject & Resubmitted, QJE*), we show how adverse selection can fundamentally limit competition by creating forces that resemble natural monopoly. Our central insight is that adverse selection strongly *intensifies price competition*, because low prices become a way of attracting or “cream-skimming” profitable healthy consumers. When multiple firms compete in this way, the result is a race-to-the-bottom in prices, which undermines the profits needed for a market to support a variety of competing firms. By linking adverse selection to the classic problem of natural monopoly, we highlight how selection can shape

market structure and why policy tools such as risk adjustment or pricing regulations are essential to sustain competition.

My working paper, “[Unobserved Heterogeneity, State Dependence, and Health Plan Choices](#)” (with Ariel Pakes, Jack Porter, and Sophie Calder-Wang; *Reject & Resubmitted, AER*) addresses insurance competition in a more indirect way. It tackles the important methodological issue of distinguishing switching costs (also called “state dependence”) from unobserved preference heterogeneity – a classic challenge that arises in many settings, including health insurance choice. We propose and implement a new moment inequalities method to estimate switching costs while allowing for extremely rich preference heterogeneity (via individual-by-product fixed effects). While this paper is less policy-centric than my other work, our empirical estimates from a health insurance setting suggest that switching costs may be smaller (and unobserved preferences larger) than previously thought.

Making competition work requires creative thinking about insurance market design. In “[Do Insurers Respond to Active Purchasing? Evidence from the Massachusetts Health Insurance Exchange](#)” (with Ethan Forsgren; *Journal of Risk and Insurance*, 2023), we study “active purchasing” policies, whereby a market regulator shapes competitive incentives to reward low-price health insurers. While entirely descriptive, this paper highlighted the large role of these policies in the Massachusetts exchange, which was able to reduce premiums by 16-20% within a two-year period after adopting them.

Ultimately, my work suggests that insurance market competition can work – but only with careful market regulation and policy incentives (like risk adjustment and subsidies) to address the pervasive issue of adverse selection. Adverse selection is not “solved” by simple mandates or subsidies, since it is deeply entwined with insurer competition and difficult-to-regulate plan features like whether they cover expensive star hospitals. My research develops new insights to help policymakers spot problems and manage the tradeoffs involved – even if it suggests insurance markets are unlikely to achieve first-best outcomes.